

Mini Report: Design of a 32-bit Multi-Cycle CPU

Weiyuan Du

1 Introduction

In this project, I implemented a 32-bit multi-cycle RISC-V CPU that supports a functional subset of instructions including R-type, I-type (ALU and Load), Store, Branch, and Jump operations. Unlike a single-cycle design, this architecture optimizes hardware usage by utilizing a shared ALU and shared memory for both instructions and data, managed through a central Finite State Machine (FSM).

2 Questions

1. **In Figure 1, some registers have write enable (*en*) while some do not. Explain why *en* is necessary or not.**

Enable signals are necessary for registers that must hold their data across multiple clock cycles, such as the PC, Instruction Register, and Register File. Without an enable, these would be overwritten every cycle. Temporary registers like ALUOut do not need enables because they only cache content until next stage.

2. **Why is the output of the sign extender not registered?**

The sign extender is a combinational circuit that takes its input from the Instruction Register (IR). Since the IR is already a register, the output of the sign extender remains stable without needing its own register.

3. **Where is the slowest datapath (critical path) between two registers, given the delays of the components are constants?**

The critical path is typically the memory access path:

PC → Address Mux → Memory → Instruction Register

4. **How does a multicycle architecture allow the clock period to be shorter than in a single-cycle one? What is the benefit?**

A multicycle design breaks instructions into smaller steps. The clock period only needs to be long enough for the longest *individual* rather than the entire instruction. This saves time by preventing simple instructions waiting all the time.

5. **Why might a multicycle design be easier (or harder) to extend with new instructions compared to a single-cycle design?**

It is easier because you can reuse existing hardware (like the ALU) by simply adding new states to the FSM. However, it is harder because the FSM logic becomes significantly more complex to design and verify as the number of instructions and states increases.

6. **Briefly describe the bug you have encountered.**

I encountered a significant bug where the control logic for Branch and R-type instructions was incorrectly encoded together. This overlap made it impossible for the ALU to stably identify the required operation, leading to frequent execution failures. I resolved this

by re-encoding the `ALUOp` signals to clearly separate branch instructions from R-type operations, ensuring the `ALUControl` logic could reliably generate the correct signals for each instruction type.

2.0.1 Performance Analysis

Instruction Type	Cycles per Instr.	# in Program	Total Cycles
R-Type (add, sub, or, etc.)	4	7	28
I-Type (addi)	4	4	16
U-Type (lui)	4	1	4
Load (lw)	5	1	5
Store (sw)	4	2	8
Branch (beq)	3	3	9
Jump (jal)	4	1	4
<i>FSM Latency / Startup</i>	-	-	2
Total	-	19	76

Table 1: Performance Analysis: Instruction Breakdown and Cycle Counts

Performance Analysis: The manual calculation for the 19 executed instructions sums to 74 cycles. The simulation reports 76 cycles ($\approx 2.7\%$ difference). This discrepancy is due to the hardware testbench capturing the initial clock cycle during the reset phase and the final cycle required for the Program Counter (PC) to update to the terminal branch address before the simulation *stop* condition is met.

Average CPI:

$$\text{CPI} = \frac{\text{Cycles}}{\text{Instructions}} = \frac{74}{19} = 3.89$$

3 System Design

3.1 FSM

Please find attached pdf page at the end of this report.

3.2 Overall Design

Please find attached pdf page at the end of this report.

4 Assembly Test Program

A complete RISC-V assembly test program was written to validate the correctness of the CPU. The test covers all implemented instruction types. Each instruction is annotated with comments explaining its expected behavior. A corresponding Verilog testbench is used to verify execution results.

Listing 1: RISC-V Test Assembly

```

1 .text
2 .global _start
3 _start:
4 lui x6, 0x00010          # x6 = 0x00010000 (65536)
5 addi x2, x0, 5            # x2 = 5
6 addi x3, x0, 12           # x3 = 12
7 addi x7, x3, -9           # x7 = (12 - 9) = 3

```

```

8 | or    x4, x7, x2          # x4 = (3 OR 5) = 7
9 | and   x5, x3, x4          # x5 = (12 AND 7) = 4
10 | add   x5, x5, x4         # x5 = (4 + 7) = 11
11 | beq   x5, x7, end        # shouldn't be taken
12 | slt   x4, x3, x4         # x4 = (12 < 7) = 0
13 | beq   x4, x0, around      # should be taken
14 | addi  x5, x0, 0           # shouldn't happen
15 | around:
16 | slt   x4, x7, x2         # x4 = (3 < 5) = 1
17 | add   x7, x4, x5         # x7 = (1 + 11) = 12
18 | sub   x7, x7, x2         # x7 = (12 - 5) = 7
19 | sw    x7, 84(x3)          # [96] = 7
20 | lw    x2, 96(x0)          # x2 = [96] = 7
21 | add   x9, x2, x5         # x9 = (7 + 11) = 18
22 | jal   x3, end            # jump to end, x3 = return address
23 | addi  x2, x0, 1           # shouldn't happen
24 | end:
25 | add   x2, x2, x9         # x2 = (7 + 18) = 25
26 | sw    x2, 0x20(x3)          # mem[104] = 25
27 | done:
28 | beq   x2, x2, done        # infinite loop

```

5 Limitation

I only checked with the assembly code inside my project. To make it more clear and readable, I didn't do a lot test bench works, except the top level one. There might be some other bugs hidden in my project.

6 Code Availability

All source code and design files are available in my GitHub repository: You can find more test bench in my repository, and also automativie tools help to reduce your work. :) <https://github.com/SinYita/riscv-single.git>



